

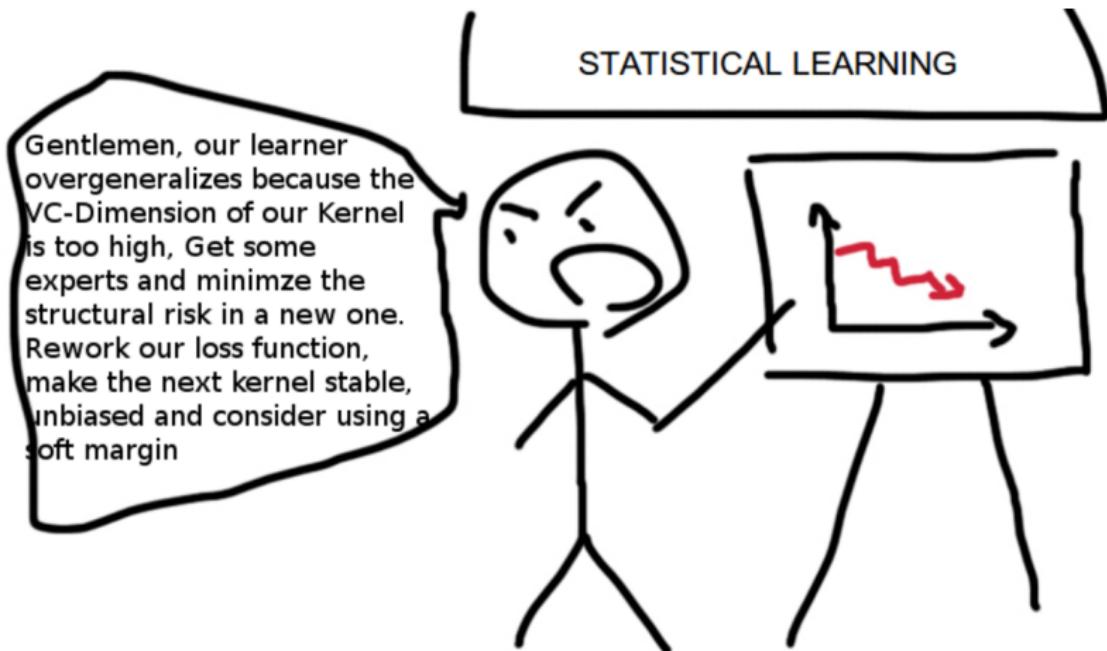
Deep Learning application to large-scale image retrieval

Pavel Nesterov

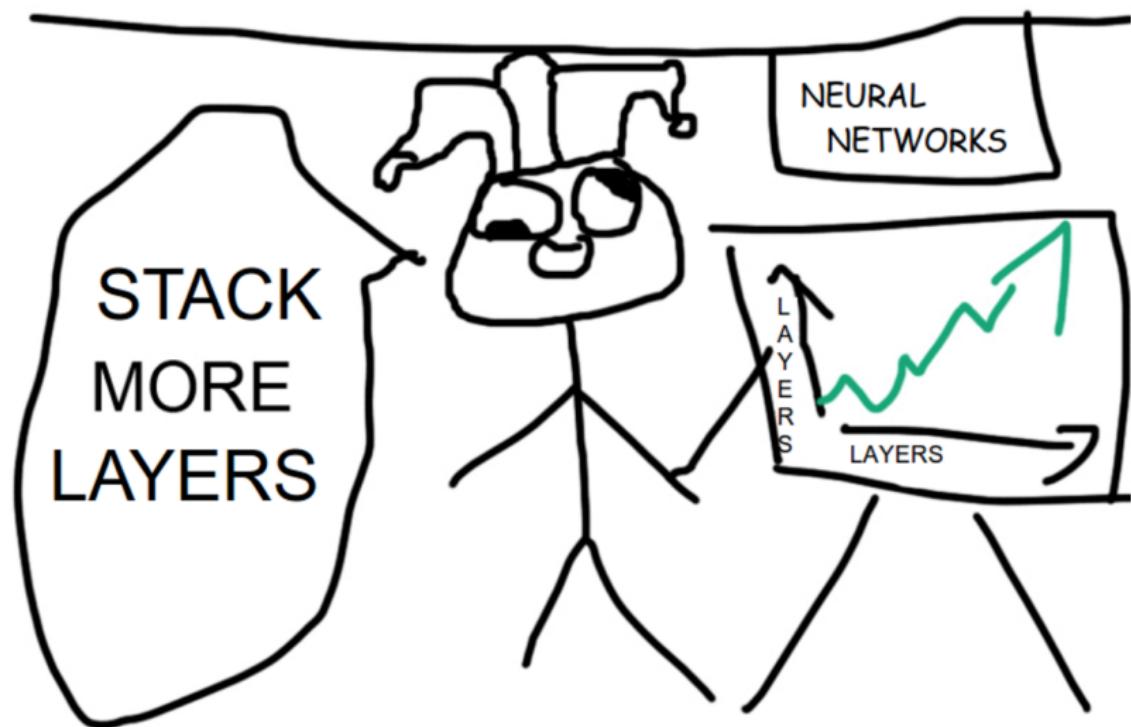
<http://pavelnesterov.info/>

November 13, 2016

Classical approach



Modern approach



Linear models

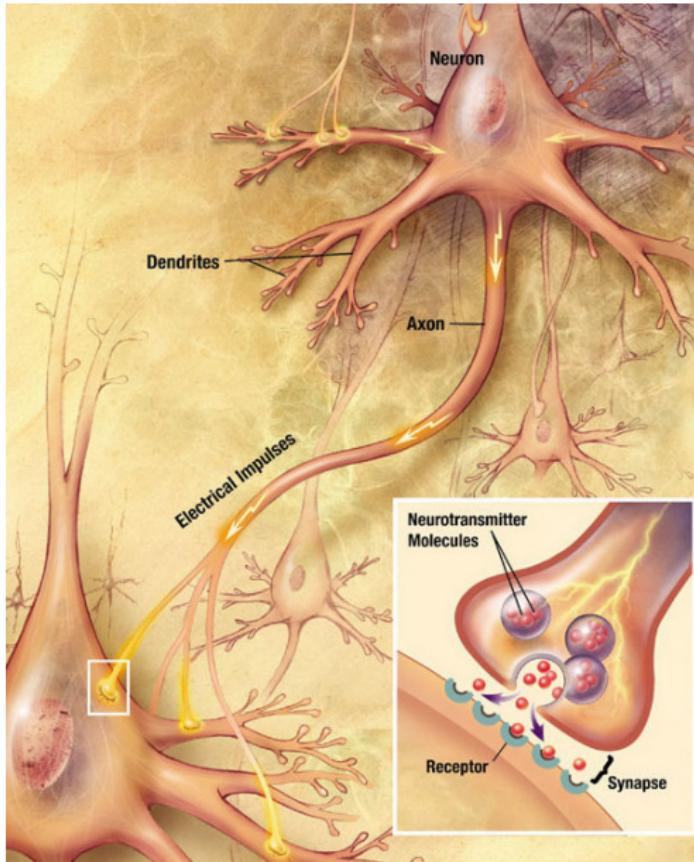
Linear regression

$$\hat{y} = w_0 + \sum_{i=1}^N w_i x_i$$

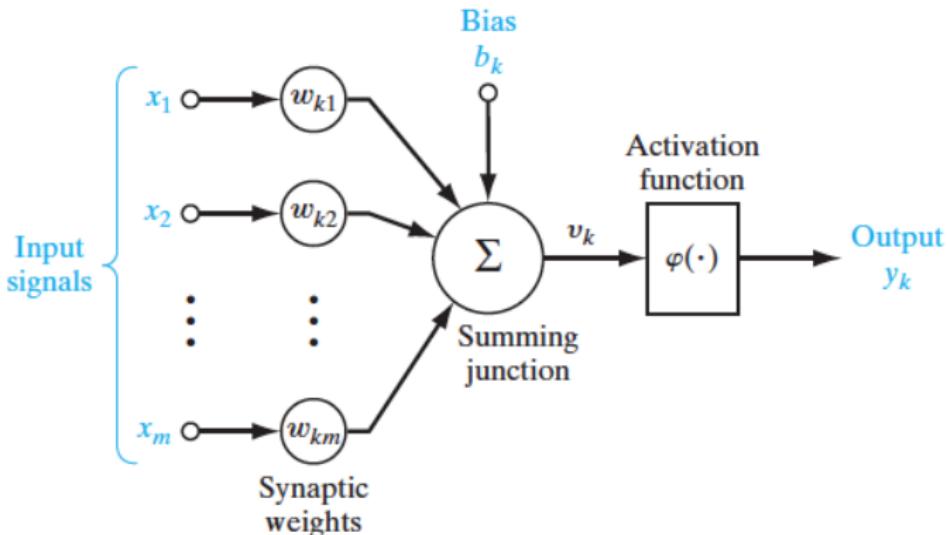
Logistic regression

$$\begin{aligned}\hat{p}(y=1) &= \sigma \left(w_0 + \sum_{i=1}^N w_i x_i \right) \\ \sigma(x) &= \frac{1}{1 + e^{-x}}\end{aligned}$$

Biological neuron



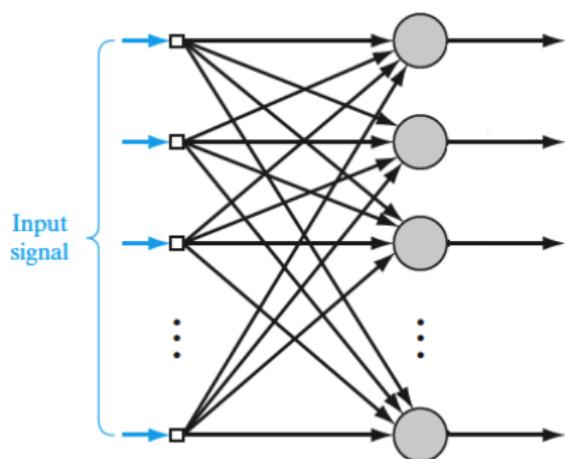
Artificial neuron



Generalized linear model

$$\hat{y} = g^{-1} \left(w_0 + \sum_{i=1}^N w_i x_i \right)$$

Single layer network



MaxEnt model

$$\begin{aligned}\hat{h}_k &= w_{k,0} + \sum_{i=1}^N w_{k,i} x_i \\ \hat{p}(y = C_k) &= \frac{e^{\hat{h}_k}}{\sum_{j=1}^K e^{\hat{h}_j}}\end{aligned}$$

Universal approximation theorem

Let we have:

- ▶ $\phi(x)$ is nonconstant, bounded, and monotonically-increasing continuous function;
- ▶ I_m is m-dimensional unit hypercube $[0, 1]^m$
- ▶ $C(I_m)$ is space of all continuous functions on I_m

then:

- ▶ $\forall f \in C(I_m) \wedge \forall \epsilon > 0$
- ▶ $\exists N \in \mathbb{N} \wedge v_i, b_i \in \mathbb{R} \wedge w_i \in \mathbb{R}^m$, where $i = 1, \dots, N$

we can define

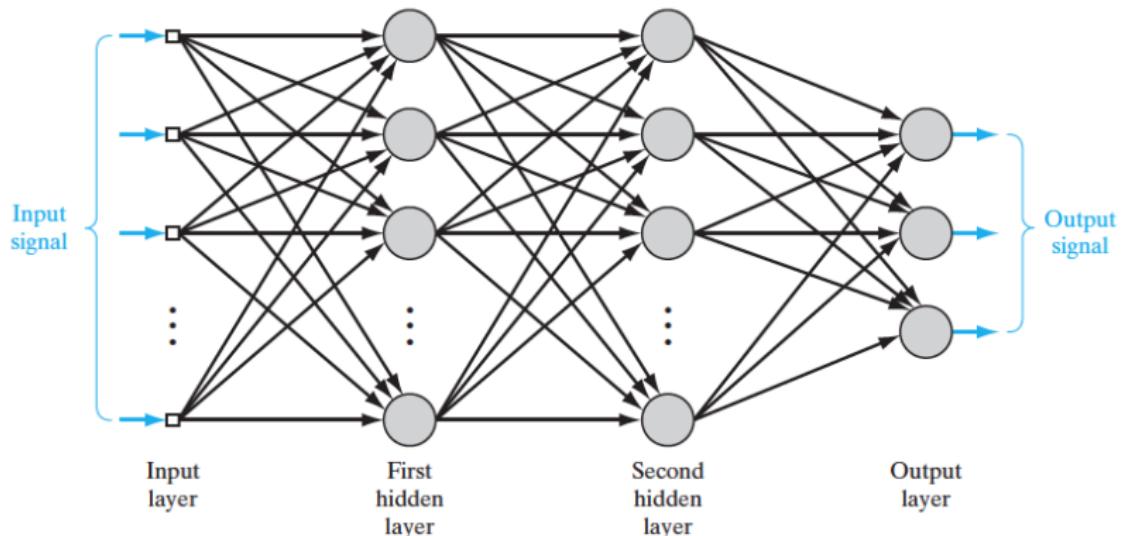
$$F(x) = \sum_{i=1}^N v_i \phi \left(b_i + \sum_{j=1}^m w_{ij} x_j \right)$$

such that

$$\forall x \in I_m : |F(x) - f(x)| < \epsilon$$

- ▶ which network topology is produced by this theorem?

Shallow network



- ▶ *how many parameters second hidden layer requires?*

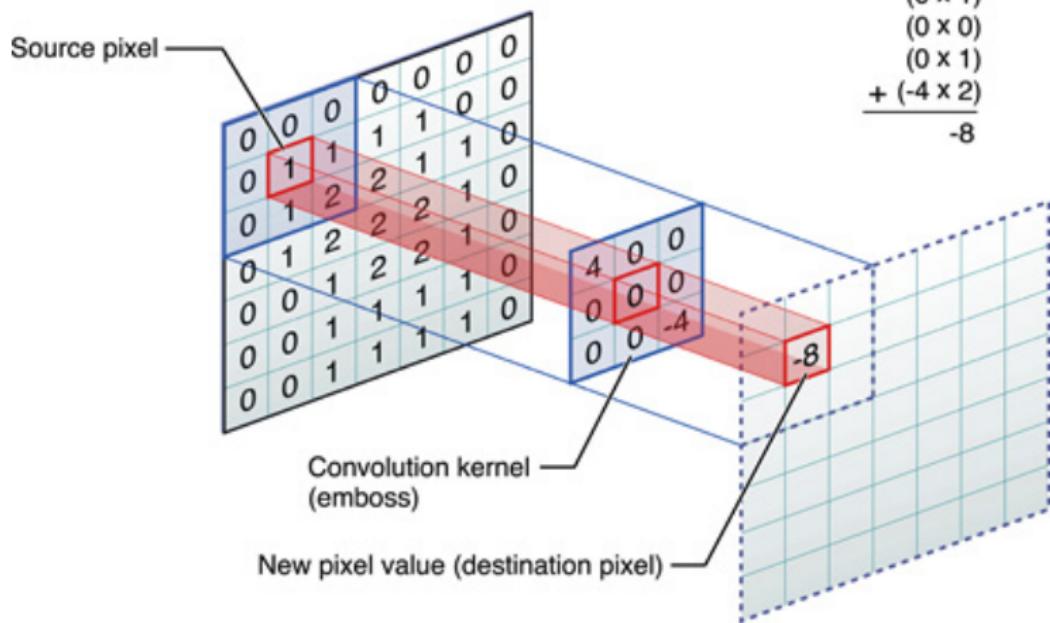
What next?



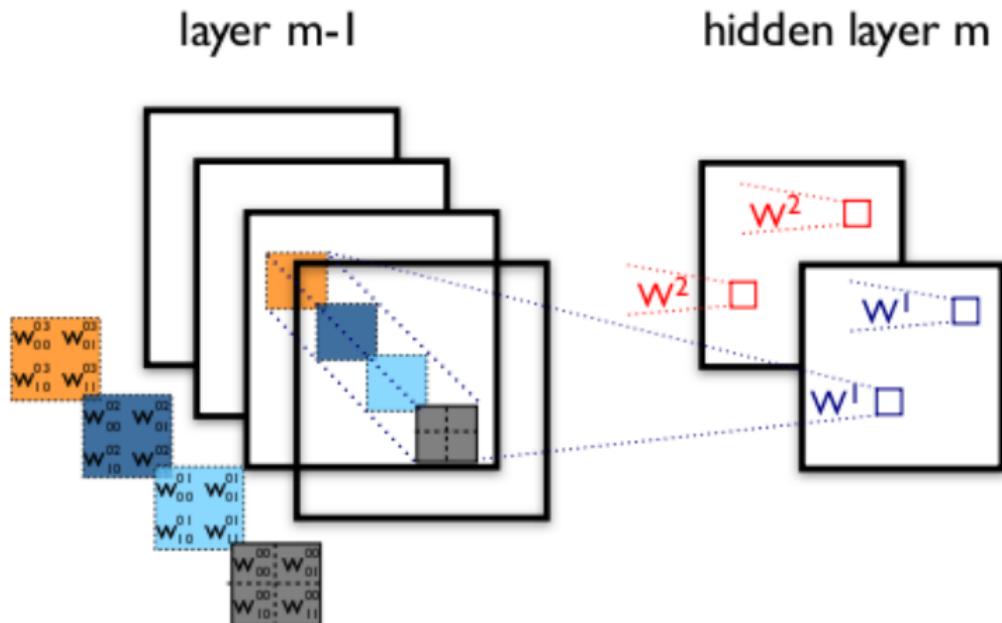
- ▶ solution is to simplify network

Convolutions

Center element of the kernel is placed over the source pixel. The source pixel is then replaced with a weighted sum of itself and nearby pixels.

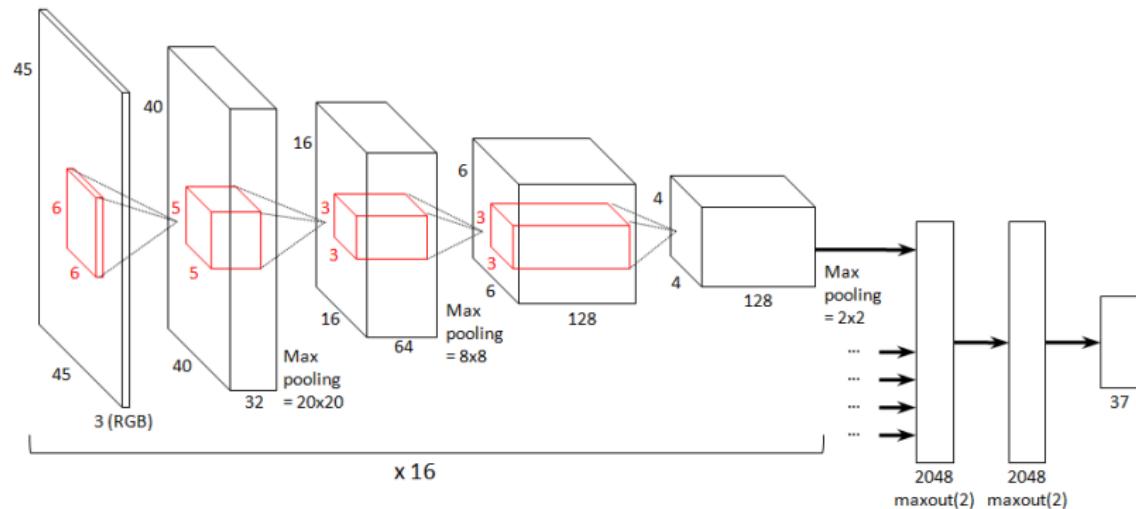


Filter bank



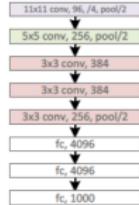
- ▶ how many parameters second convolutional hidden layer requires now?

Deep convolutional network

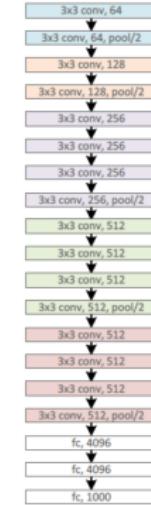


Very deep convolutional network

AlexNet, 8 layers
(ILSVRC 2012)



VGG, 19 layers
(ILSVRC 2014)



GoogleNet, 22 layers
(ILSVRC 2014)



Very very deep convolutional network

AlexNet, 8 layers
(ILSVRC 2012)



VGG, 19 layers
(ILSVRC 2014)



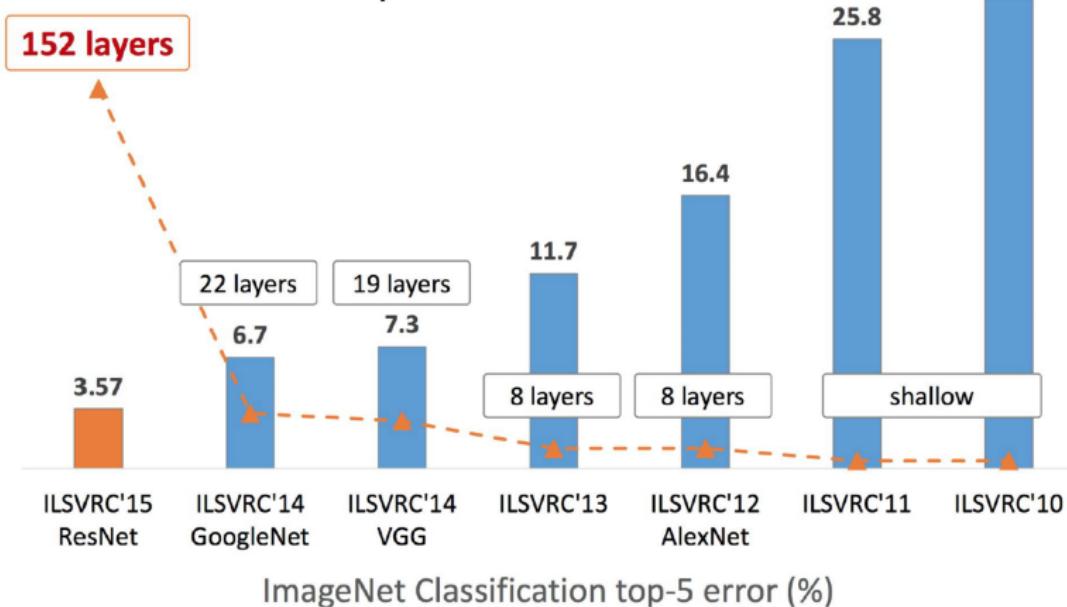
ResNet, 152 layers
(ILSVRC 2015)



MOAR LAYERS

MOAR MOAR MOAR

Revolution of Depth



Opening the black box

Given a network F with parameters Θ and input image x , lets also define output of each layer as F_i :

Training

Understanding

$$\Delta\Theta = -\lambda \frac{\partial F(x, \Theta)}{\partial \Theta}$$

$$\Delta x = \frac{\partial F_i(x, \Theta)}{\partial x}$$

How network see the world, #1

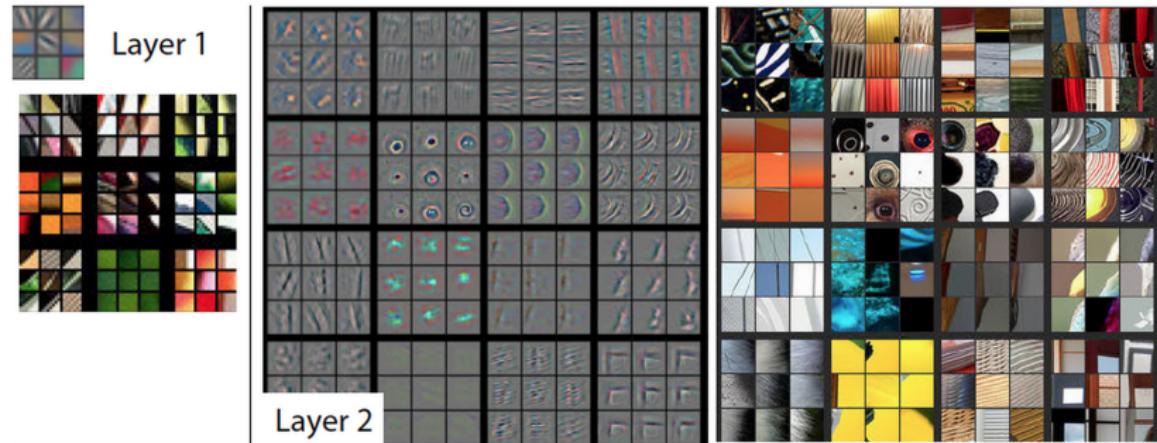
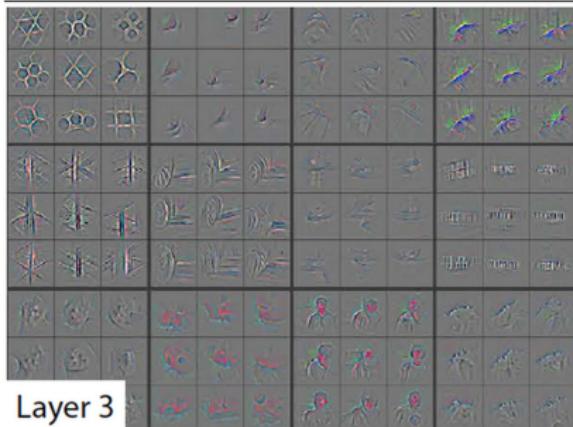


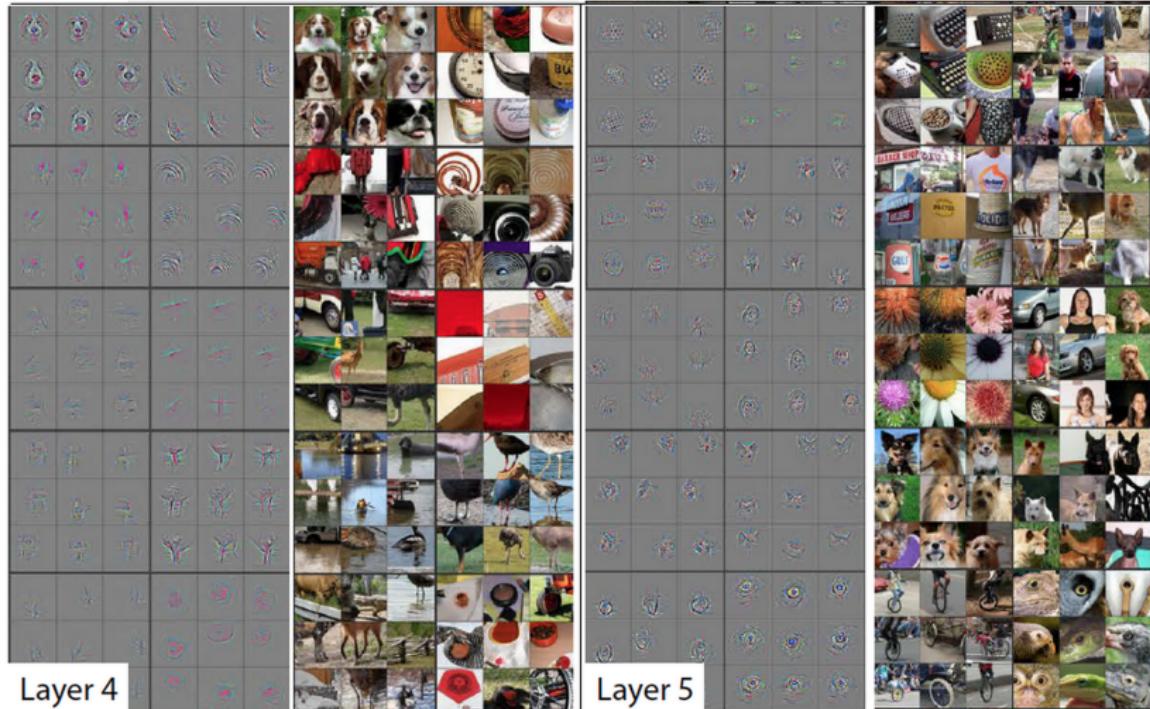
Figure 1: Visualizing and Understanding Convolutional Network ¹

¹Matthew D. Zeiler and Rob Fergus

How network see the world, #2



How network see the world, #3



Layer 4

Layer 5

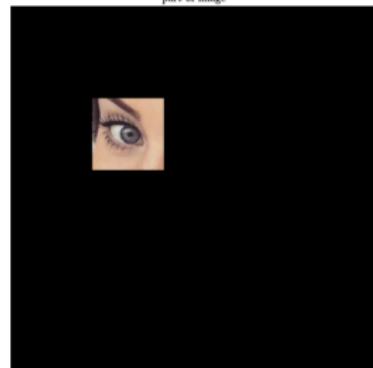
Eye filter, #1

conv4'3: max(filter=0)

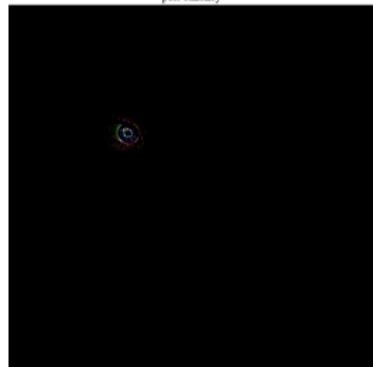
input



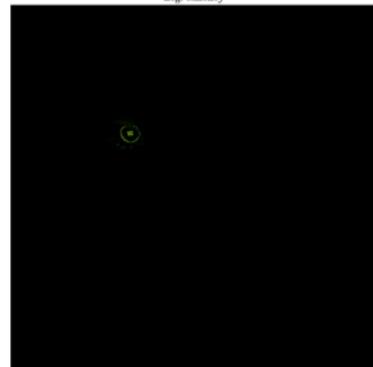
part of image



pos. saliency



neg. saliency



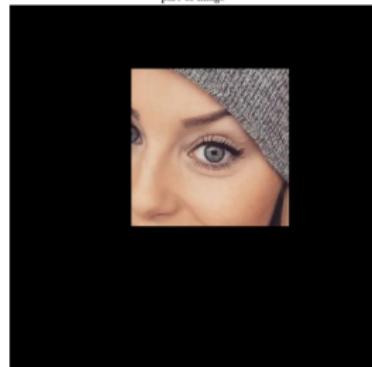
Eye filter, #2

conv5'3; max(filter=0)

input



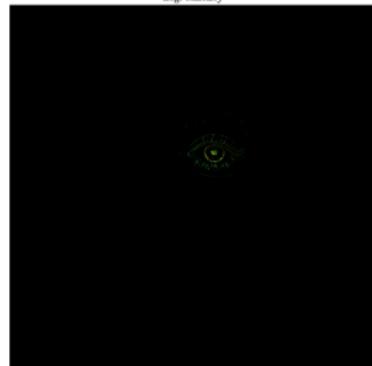
part of image



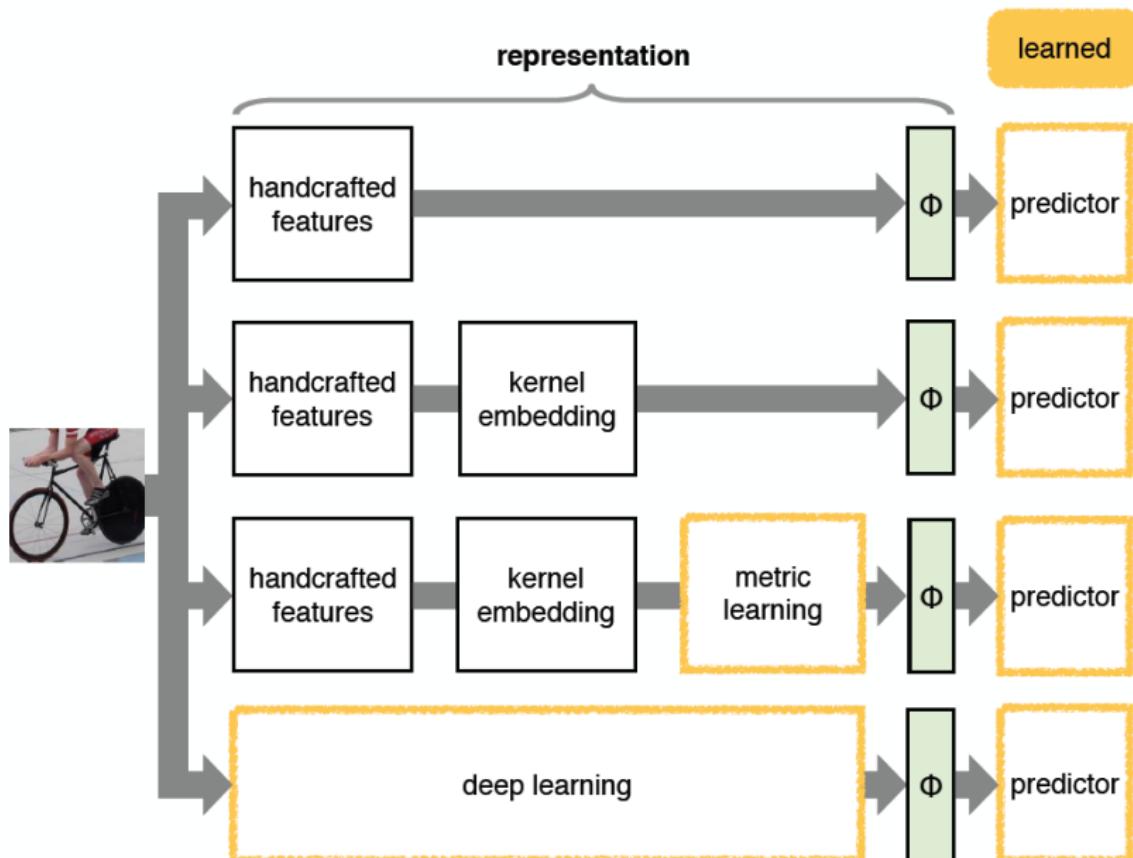
pos. saliency



neg. saliency



Deep learning = Representation learning



Semantic hashing

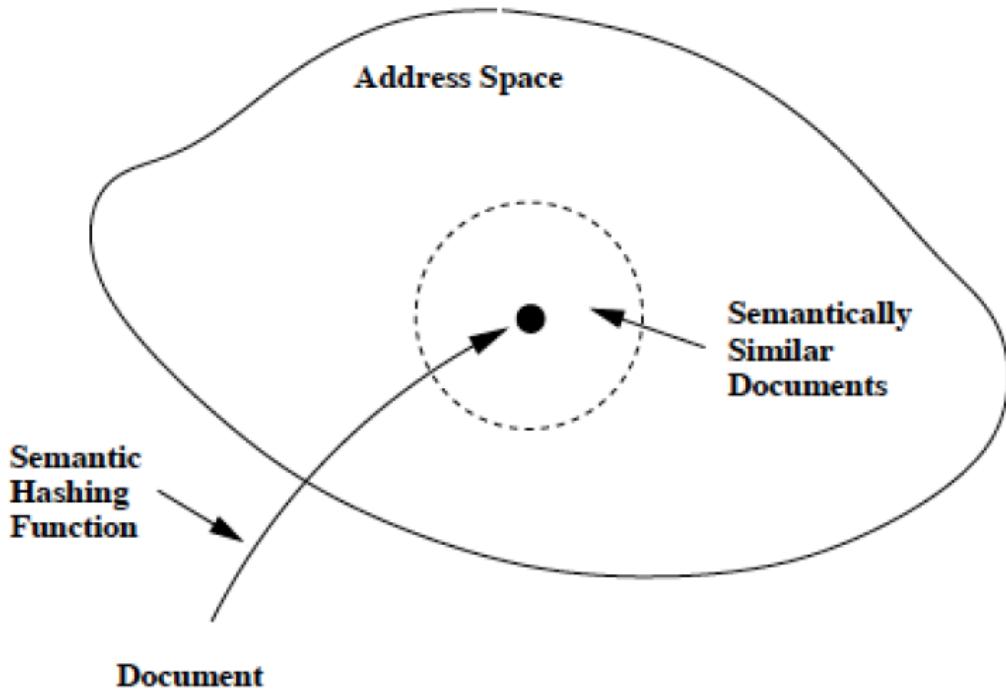


Figure 2: Semantic representation²

²Semantic Hashing (Salakhutdinov, Hinton)

Deep sparse autoencoder

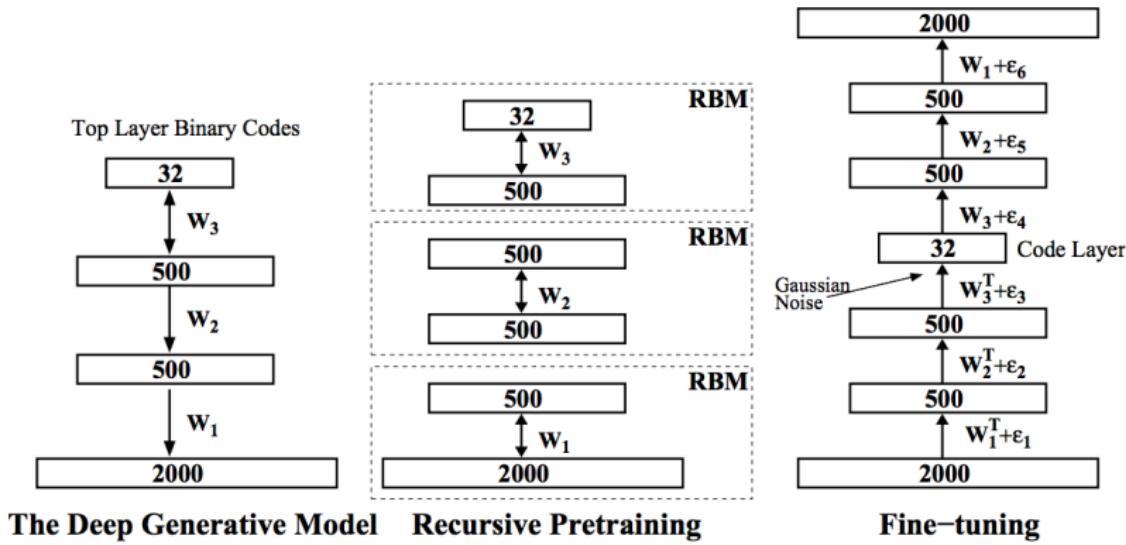
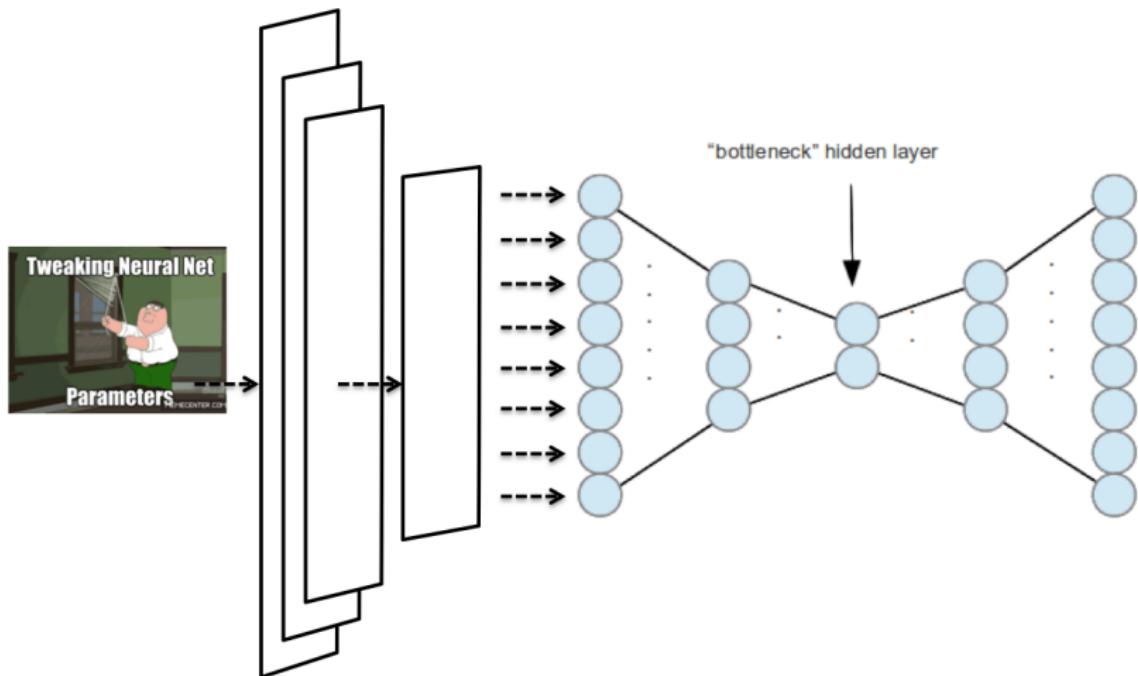


Figure 3: Training of deep autoencoder³

³Semantic Hashing (Salakhutdinov, Hinton)

Compression



Visualization, #1



1550825.jpg



1620216.jpg



1700151.jpg



1700178.jpg



1701017.jpg



1764506.jpg



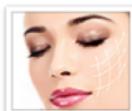
1764518.jpg



1877486.jpg



2026084.jpg



2094066.jpg



2109108.jpg



2109153.jpg

Visualization, #2



0goal.jpg



93186.png



171581.png



364466.jpg



629451.jpg



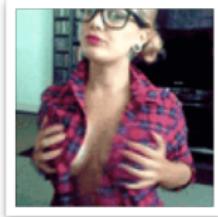
633182.jpg



712078.jpg



808167.png



814004.png



1483664.jpg

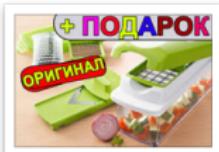


1622336.jpg

Visualization, #3



0goal.jpg



346763.jpg



362624.jpg



444028.jpg



513224.jpg



534278.jpg



571146.jpg



768177.png



847001.jpg



867038.jpg



1604230.jpg

Visualization, #4



0goal.jpg



378829.jpg



611505.jpg



689011.jpg



899429.jpg



930662.jpg



1020816.jpg



1149542.jpg



1272985.jpg



1282751.jpg



2148107.jpg

Visualization, #5

2276683425790278213



1115802.jpg



1115868.jpg



1791723.jpg



1897651.jpg



1897678.jpg



2006720.jpg



2006725.jpg



2006727.jpg



2078335.jpg



2084644.jpg



2084941.jpg



2102445.jpg

Visualization, #6



1734142.jpg



1734143.jpg



1734153.jpg



1734156.jpg



1734164.jpg



1734165.jpg



1961146.jpg



2014413.jpg



2014426.jpg



2065165.jpg



2069372.jpg



2069437.jpg



2082477.jpg



2100696.jpg



2105538.jpg

Visualization, #7



0goal.jpg



686049.jpg



704389.jpg



704390.jpg



795077.jpg



950742.jpg



1054508.png



1179832.jpg



1367326.jpg



1429950.jpg



1628913.jpg